

User Guide
mARGOt framework

Draft version

November 2, 2017

Contents

1	Introduction	1
1.1	Target problem	1
1.2	Framework overview	1
1.3	Framework integration	3
2	Monitors module	4
2.1	Application Goals	6

1 Introduction

This document describes the problem addressed by mARGOt and it shows how the framework aims at solving it, highlighting the boundaries between what is expected as input and what is the produced output. Moreover, this guide will describe the framework interface exposed to the user and how to integrate mARGOt in the target application. This document aims at providing the “big picture” of the framework, for implementation details refers to doxygen documentation.

Together with the autotuning framework, it is possible to use an utility, named *mARGOt heel*, which generates an high level interface, hiding as much as possible the implementation details of the integration process. However, the latter is considered out of scope from this document.

1.1 Target problem

The source code describes the functional behavior of the application and it might be abstracted as a function that produces the desired output, starting from an initial input, i.e. $o = f(i)$. A large class of application might expose software knobs which alter its behavior, i.e. $o = f(i, k_1, k_2, \dots, k_n)$. The idea is that a change on the configuration of those knobs, alter the extra-functional properties (EFP) of the function f (e.g. the execution time or the power consumption) and of the output (e.g. result accuracy or size).

While the functional behavior of the application is utterly described in the source code, its extra-functional behavior is more subtle, since it usually depends on the underlying architecture, on the system workload, on the assigned computational resources and on the current input. Since the user typically has requirement on the application EFP, selecting the most suitable configuration of the software knobs is not a trivial task. Moreover, since the system workload, the assigned resources and the input may change at runtime, choosing a one-fits-all configuration may lead to a sub-optimal solution.

In the context of autonomic computing, an application is seen as an autonomous element, which is able to automatically adapt. The mARGOt autotuning framework aims at enhancing the target application with a runtime adaptation layer, to continuously provide the most suitable configuration according to application requirements and to observation of its actual EFP.

1.2 Framework overview

The main idea behind mARGOt is the MAPE¹ loop. In this context, an autonomic manager is in charge to manage the application, implementing a loop composed by five elements. The **M**onitor element provides the ability to gather insight on the actual application behavior, the **A**nalyse element extract information useful for the **P**lan element, which select the action actuated by the **E**xecute element. The effect of the selected action will be observed by the monitor element, which close the loop. The fifth element of the MAPE loop is the application knowledge, which is leveraged by all the other elements.

¹Kephart, Jeffrey O., and David M. Chess. “The vision of autonomic computing.” *Computer* 36.1 (2003)

The autotuning framework is composed by three modules. Obviously, the Monitors module is in charge of acquiring insight on the actual behavior of the application, representing the monitor element of the MAPE loop. The Manager module represents the analyze and plan elements of the MAPE loop, to provide to the application the most suitable configuration according to the monitor information, the application knowledge, the features of the current input and the user requirements. The application knowledge is an abstraction of the extra-functional behavior of the application, gathered at design-time through a Design Space Exploration. However, it is possible to change the application knowledge at runtime.

1.3 Framework integration

The purpose of this document is to describe how it is possible to integrate the mARGOt framework in the target application. Following sections of this document will describe in details the interface and purpose of the monitors, the manger and the application knowledge. Finally, the section will provide an example of integration in a toy application.

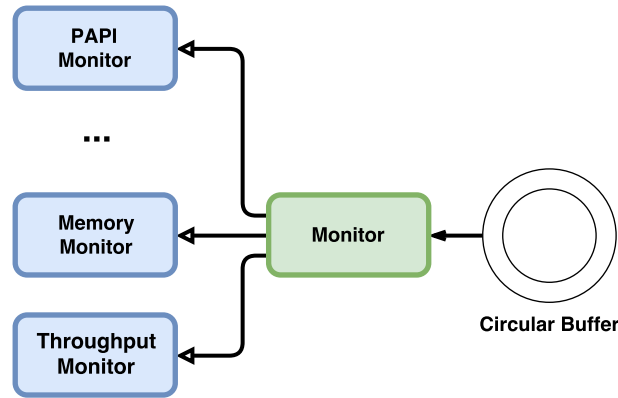


Figure 2: Overview of the monitors module.

2 Monitors module

The most important step in order to adapt, is to acquire insight on the current behavior of the application EFP of interest, which defines the performance of the application. For example, the application developer might be interested on the execution time of the kernel or in the accuracy of the result. While the first metric might be applied to every class of applications, the accuracy of the result is meaningful only in the context of an application. For example, a video encoder might use the Peak to Signal Noise Ration (PSNR) as accuracy metric, while a scientific application might be more interested on the difference between the actual result and the ground truth value.

For this reason the monitors module must be flexible enough to let the user to define custom monitor to observe application-specific metrics. Figure 2 shows the overview of the structure of the monitors module. A base class (*Monitor*) implements a circular buffer which store the last n (user defined) observations and provides to the user the possibility to extract statistical properties over observed data. Moreover, this class implements all the methods required to interact with all the other framework elements. It is possible to specify the type of the elements stored in the circular buffer and the precision used to compute the statistical properties (by default it uses at least 32bit float). Each time a statistical properties is computed, the monitor uses a cache-like mechanism to avoid to recompute the property.

The actual monitor must extend the latter class, with the functionality that actually gather the target metric and push the value in the circular buffer (using the *push*) method. mARGOt ships with the following suite of monitors:

- Energy monitor (uses the RAPL environment)
- Frequency monitor (uses the CPUfreq environment)
- Memory monitor
- Odroid power monitor

```

1  // monitor definitions
2  margot::Monitor<float> my_float_monitor(3);
3  margot::Monitor<int, double> my_int_monitor;
4  margot::Monitor<double> my_double_monitor;
5  margot::TimeMonitor timer;
6
7  // how to observe a metric
8  timer.start();
9  timer.stop();
10 my_float_monitor.push(3.2f);
11
12 // how to extract statistical properties
13 const auto avg_time = timer.average();
14 const auto avg_float = my_float_monitor.average();

```

Figure 3: Example of C++ code to define and use monitors.

- Odroid energy monitor
- PAPI monitor (uses the PAPI environment)
- Process and system CPU usage
- Temperature monitor (uses the Sensors environment)
- Throughput monitor
- Time monitor

Figure 3 shows some examples on how it is possible to define and use monitors. In particular, line 1 defines a basic monitor which stores float and has a circular buffer of three elements. Line 2 defines a monitor of integer, however it specifies that the statistical properties such as average and standard deviation should be performed using 64bit of precision. Moreover, the default constructor sets the circular buffer size to just one element. Line 3 defines a monitor of doubles. Since the default type used to compute statistical properties is a float, mARGOT automatically promotes the latter to double, to prevent precision loss. Line 4 defines a time monitor, which is a specialization of the generic monitor, suitable to record the elapsed time from a start and stop method (lines 8,9). In general each specialized monitor provides its own method to observe the target metric. For a full description of each monitor, refer to the doxygen documentation. To “observe” a value using a custom monitor (line 10), it is required to use the *push* method. Regardless of the monitor type, to extract statistical properties (lines 13-14), it is enough to use the dedicated method. In particular, it is possible to extract the following statistical properties (defined in the enum *DataFunctions*):

- Average
- Standard deviation

```

1   margot::TimeMonitor timer(margot::TimeUnit::MILLISECONDS);
2   margot::Goal<int, margot::ComparisonFunctions::LESS> my_goal(2);
3
4   timer.start();
5
6   // application kernel
7
8   timer.stop();
9
10  if (my_goal.check<float, float, margot::DataFunctions::AVERAGE>(timer))
11  {
12      std::cout << "We are fast" << std::endl;
13  }
14  else
15  {
16      std::cout << "We should improve" << std::endl;
17  }

```

Figure 4: Example of C++ code to check if the elapsed time to execute the application kernel is below $2ms$.

- Maximum element observed
- Minimum element observed

2.1 Application Goals

Since the application developer might have requirements on the lower bound or upper bound on a metric of interest. mARGOt uses the class *Goal* to represent this concept. In particular, it is possible to test if a statistical properties of the monitor achieve the goal or to compute its absolute or relative error. In the current implementation, it supports four standard comparison functions (greater than, greater or equal than, less than and less or equal than), defined in the enum *ComparisonFunctions*.

Figure 4 shows an example on how it is possible to use the goals and monitors, to check if the application kernel execution time is below $2ms$. In particular, line 1 defines a time monitor with a resolution of milliseconds. Line 2 defines a goal with value two and comparison function “less than”. After that we have instrumented the kernel of the application (lines 4-8), we check if the execution time was below $2ms$ (line 10), and we print an output message accordingly (lines 10-17).